

The Bochner loss closes the variance gap across architectures (plain L^2 collapses it; the Bochner loss restores var. ratio $\rightarrow 1$)

